

Social networks of lexical innovation

Investigating the diffusion of neologisms on Twitter

Quirin Würschinger
LMU Munich

16th June 2020

Todo list

page?	5
freq. esp. insufficient for lex. inn.	9

Abstract

Societies continually evolve and speakers coin and use new words to talk about innovative products and practices. While most lexical innovations fail to catch on, others spread successfully and become part of the lexicon. This paper investigates the diffusion of English neologisms on Twitter. Previous work on lexical innovation has almost exclusively relied on usage frequency counts for measuring diffusion. Taking frequency as a baseline, we use social network analysis to zoom in on the sociolinguistic dynamics of diffusion.

Our results show that frequency counts lend themselves to approximate overall degrees of diffusion with varying success. While absolute counts can be misleading, incorporating temporal dynamics of use provides a better picture of diffusion. However, frequency-based information alone fail to capture important sociolinguistic characteristics. Social network information are shown to add valuable information about whether new words are known and used by an increasing number of individuals and communities of speakers. Firstly, we distinguish different pathways of diffusion depending on whether and to which degree new words show increasing vs. decreasing centralized use over time. Secondly, we show that social network information allow for more fine-grained assessments of degrees of diffusion, for example when new words are used with increasing frequency while their remains limited to certain parts of the speech community. Lastly, we compare our results based on usage frequency and on social network analysis. Besides notable discrepancies, we find a significant correlation between both types of information which serves to cross-validate both approaches.

Our results suggest that social network information can complement frequency counts and that using information from both sources provides a more reliable and differentiated view of the sociolinguistic dynamics of diffusion. We suggest that this is particularly important for investigating the diffusion of lexical innovations, as new words are often marked by high social indexicality and show substantial differences in use between communities of speakers. More generally, however, social network analysis shows great potential to study sociolinguistic dynamics of language variation and change on all linguistic levels.

Keywords: lexicology, lexical innovation, sociolinguistics, diffusion, social media, Twitter, big data, social network analysis

1 Introduction

Societies continually evolve, new products and practices emerge, and speakers coin and adopt new words when they interact and share information. How do these new words spread in social networks of communicative interaction?

Covid-19 has recently spread through social contagion with shocking speed and has tragically affected the lives of people around the world. Its fatal consequences have demonstrated the devastating power of exponential diffusion in social networks. In a recent paper analysing contagion patterns of diseases in *Nature Physics*, Hébert-Dufresne, Scarpino & Young (2020) suggested that the spread of viruses follows principles of complex contagion through social reinforcement and that it matches the dynamics of diffusion

of cultural and linguistic innovations such as new words and internet memes. Does this confirm the widespread perception that new words ‘go viral’?

Influential sociolinguistic models of the spread of linguistic innovations like the S-curve model (Milroy 1992) share fundamental features with earlier economic models (Rogers 1962) of diffusion and show commonalities between the spread of cultural and linguistic innovations. These models assume that diffusion in social networks follows universal trajectories and that rates of spread depend on social dynamics such as network density and the presence or absence of weak ties (Granovetter 1977). Unlike research on biological and cultural diffusion processes, however, sociolinguistic research has only recently been provided with data sources that are equally suitable for large-scale, data-based approaches using social network analysis to study these phenomena empirically.

Social media platforms like Twitter have changed the way we communicate and how information spreads, and they offer large amounts of data for empirical research. Sociological research has been concerned with pressing issues regarding the impact of online social networks for the spread of hate speech, fake news and the power of ‘influencers’, bots and institutions on public opinions and elections, which increasingly strain the social fabric. For (socio-)linguists, social media provide large amounts of data of authentic language use which opens up new possibilities for the empirical study of language variation and change. The size of these datasets as well as their informal nature allow for large-scale studies on the use and spread of new words, for example, to gain insights about general trajectories (Nini et al. 2017) or about factors that influence whether new words spread successfully (Grieve 2018). Moreover, metadata about speakers allows studying aspects of diffusion that go beyond what can be captured by usage frequency alone. Recent work has, for example, used Twitter data to investigate the geographical spread of lexical innovations. (Eisenstein et al. 2014, Grieve 2017, Grieve, Nini & Guo 2018)

Data about the communicative interaction of speakers additionally allows performing network analyses of the social dynamics of diffusion processes. Network science approaches to social media data have been successfully employed in diverse fields, for example, to study the spread of diseases (Lu et al. 2018), opinions (West et al. 2014) and political attitudes (Pew Research Center 2019). While the study of social networks has a long research tradition in sociolinguistics and has shaped influential models of diffusion (e.g. Milroy & Milroy 1985)), large-scale network analyses of sociolinguistic phenomena have only recently become more widespread. These new data sources and methodological advances put computational sociolinguistics in an excellent position to gain new insights and to test long-standing theoretical models empirically.

In the area of lexical innovation, this can serve to evaluate important theoretical concepts like the role of early adopters, network density and weak ties in the diffusion of new words. For example, earlier approaches have used computational modelling to test the validity of the S-curve model (Blythe & Croft 2012), and to model processes of simple and complex contagion of linguistic innovations in social networks (Goel et al. 2016). Applying social network analysis to bigger samples of neologisms and tracking their diffusion on social media datasets promises to shed light on whether the adoption of new words remains limited to closely connected sub-communities or whether they

reach larger parts of the speech community and whether individuals or groups drive this process.

This paper makes use of social media data and social network analysis to study the diffusion of lexical innovations on Twitter. Taking usage frequency as a baseline, we conduct a longitudinal study monitoring the use of a broad sample of neologisms to analyse their cumulated usage frequency as well as the temporal dynamics underlying their spread. We additionally conduct social network analyses of our neologism sample to get a better picture of the sociolinguistic dynamics at play to assess different pathways and overall degrees of diffusion. Lastly, we compare both approaches to assess their validity, and we combine information from both sources to draw a more differentiated picture of diffusion.

The paper is structured as follows. Section 2 presents an overview of previous attempts to modelling and measuring the diffusion of lexical innovations in order to contextualise and define the present theoretical framework and its operationalisation for the empirical study. Section 3 provides information regarding the present sample of neologisms and the collection and composition of the Twitter dataset this study is based on. Section 4 describes the methodological procedure for analysing diffusion in this dataset, focusing on the construction and analysis of social networks. Section 5 presents the empirical results obtained from using usage frequency and social network analysis to study diffusion, and from comparing both approaches. Section 6 summarises and discusses these results and suggests theoretical implications and directions for future work.

2 Modeling the diffusion of lexical innovations

Speakers continually coin new words, yet most fail to spread successfully and fall into oblivion. How do new words diffuse to be known and used by more and more speakers and to become conventional lexemes in a language system? And how can diffusion be modelled theoretically and measured empirically?

Neologisms are on a continuum from entirely novel word-formations to established lexemes that are familiar to the majority of the speech community. Neologisms have spread to some extent, but are still perceived as new or unknown by many speakers. On one end of the continuum, ‘nonce-formations’ are new words that have been coined in a concrete communicative situation, but are not adopted by interlocutors to be used in future usage contexts and do not enter a process of continuous diffusion. (Hohenhaus 1996) Fully established words form the other end of the continuum. They are known and used by most or all members of the speech community and codified in dictionaries. This latter, lexicographic feature marks speakers’ agreement on how these words are to be used and is a sign of their status as conventional lexical units in the language system.

Neologisms occupy an intermediate position between both poles and can be defined as

[...] lexical units, that have been manifested in use and thus are no longer nonce-formations, but have not yet occurred frequently and are not widespread enough in a given period to have become part and parcel of the lex-

icon of the speech community and the majority of its members. (Kerremans 2015: 31)

Diffusion thus represents the process that transports successful neologisms along this continuum, becoming increasingly conventional in the speech community.

A more precise definition is provided by Schmid: 'I define diffusion as a process that brings about a change in the number of speakers and communities who conform to a regularity of co-semiotic behaviour and a change in the types of cotexts and contexts in which they conform to it.' (Schmid 2020)

this definition includes three dimensions of diffusion

cotext context speakers

i will focus on sociolinguistic dimension in this paper spread across speakers communities

previous work has taken at structural and cognitive perspectives

i will focus on the sociolinguistic perspective

most relevant and important model: S-curve model

page?

2.1 Research perspectives

A substantial body of linguistic research has tackled this question from different **perspectives**. (Schmid 2016: 16)

From a **structural** perspective, main areas of interest include which word-formation processes are involved in forming new words, whether they are formally and semantically transparent, whether they change in the process of lexicalization and which status the resulting neologisms have in the language system (institutionalization). (e.g. Bauer 1983, Lipka 2005)

Cognitive perspectives focus on how individuals process and store lexical innovations. Speakers generally use new words when they experience a communicative need to talk about entities or practices that cannot be expressed by their language's inventory of conventional words yet. In order for neologisms to successfully diffuse, speakers need to successfully negotiate their meaning (co-semiosis) in discourse, others need to adopt the behaviour of using these words (co-adaption). Continued exposure and use of new words can then lead to the entrenchment of new words in the mental lexicon of speakers. (Schmid 2008)

Sociolinguistic perspectives transcend the level of the individual to study the diffusion of new words across speakers. The diffusion of lexical innovations is commonly seen as successful when the majority of the speech community has accepted a new word as a conventional lexical unit that can and is being used in communicative practice.

2.2 The S-curve model

S-curve models of linguistic change (Milroy 1992, Nevalainen 2015, Labov 2007) assume universal sociolinguistic dynamics for the diffusion of linguistic innovations.

- The **trajectory** of spread is expected to follow an S-curve shape, with low rates of diffusion in early stages, followed by a period of accelerating spread with a tipping

point at the mid point in the diffusion curve after which diffusion slows down and the curve flattens towards the end of the diffusion process.

- These temporal trajectories are assumed to correspond to the **sociolinguistic dynamics** of which individuals and groups interact with each other and adopt the target innovation.
 - In the **first stage** of slow diffusion only few early adopters take up the innovative words. The individuals who use the new word typically form dense networks connected by strong ties. The structure of tight-knit communities of potentially like-minded individuals of similar backgrounds facilitates the successful negotiation of meaning (co-semiosis) of new words. High rates of interactions in these communities leads to high rates of exposure for individuals, which fosters co-adaptation, entrenchment and usualization of new words in these communities.
 - In cases of successful diffusion the initial stages are followed by an **acceleration in spread** when new words increasingly reach speakers outside these tight-knit communities via weak ties (Granovetter 1977). Rates of diffusion increase substantially when speakers that are not part of the initial group of early adopters start to accommodate the new words, allowing the innovations to reach a broader spectrum of the speech community.
 - In **later stages**, rates of diffusion slow down again as the majority of the speech community has already adopted the new words while smaller pockets speakers remain reluctant to take up the new words.
- S-curve models have mainly been applied to the **linguistic domains** of phonology and syntax. Fundamental differences between lexemes and linguistic items on other levels such as phonemes and grammatical constructions might affect the validity and reliability of such models for *lexical* innovation.
 - For example, grammatical constructions such as the *going to* future used to express a speaker's future intention serve to fulfil relatively abstract communicative needs that remain stable over time. By contrast, on the lexical level, linguistic innovations are typically tied to concrete cultural referents such as products and practices whose conceptual relevance is much more volatile over time. For example, many lexical innovations such as *millennium bug* denoting the fear of a computer crash at the beginning of the new millennium can show high rates of diffusion and become entrenched and conventional among the majority of the speech community. Without continual conceptual relevance in public discourse, however, these words fail pass on to the next generation of speakers. S-curves are commonly assumed to be found when linguistic innovations compete for '**semantic carrying capacity**' (Nini et al. 2017), however, in many if not most cases of *lexical* innovation the conceptual carrying capacity is far from stable over time which represents a critical deviation from the traditional assumptions behind S-curves in language change.

- However, the generally strong theoretical and empirical basis of the S-curve model for language innovation and change, also from studies on the diffusion of cultural innovations (Rogers 1962), and the precise formulation of the sociolinguistic dynamics underlying different phases of diffusion still make it an attractive **blueprint** for the empirical study of the sociolinguistic diffusion of lexical innovations.

2.3 Current framework: the EC-Model (Schmid 2020)

I use the Entrenchment-and-Conventionalization-Model (Schmid 2020) as a framework for modelling the diffusion of lexical innovations.

The EC-Model provides an approach integrating both structural, cognitive and sociolinguistic perspectives on the diffusion of lexical innovations.

The model also differentiates between the level of the individual ('entrenchment') and the community ('conventionalization').

Here I will only briefly outline the most important concept relevant for studying the sociolinguistic aspects of diffusion here.

Conventionalization:

definition: 'Conventionalization is the continual process of establishing and re-adapting regularities of communicative behaviour among the members of a speech community, which is achieved by repeated usage activities in usage events and subject to the exigencies of the entrenchment processes taking place in the minds of speakers.' (Schmid 2020)

Usualization 'Usualization can therefore be defined as a process that establishes, sustains, and changes regularities of behaviour with regard to co-semiotic mappings between forms and meanings or functions and communicative goals and linguistic forms. It affects the semasiological, onomasiological, syntagmatic, cotextual, and contextual dimensions of conformity behind conventionality and is relative to communities.' (Schmid 2020)

Diffusion 'Linking the three aspects of speakers, cotexts, and contexts, I define diffusion as a process that brings about a change in the number of speakers and communities who conform to a regularity of co-semiotic behaviour and a change in the types of cotexts and contexts in which they conform to it.' (Schmid 2020)

less relevant for sociolinguistic aspects and for this study cotexts contexts

According to the EC-Model, for studying the sociolinguistic aspects of diffusion, investigating 'changes in the number of speakers and communities' is thus essential.

3 Measuring the diffusion of lexical innovations

3.1 Previous approaches

- **before frequency** Empirical approaches studying the diffusion of lexical innovations have only recently become feasible with the advent of new data sources and computational methods.

- Earlier work had to rely on **traditional linguistic corpora**. Due to the low-frequency nature of neologisms general linguistic corpora do not allow studying broad ranges of neologisms and thus pose limits to making broad and robust generalizations about the nature lexical innovation. Despite these limitations, case studies on selected neologisms (Hohenhaus 2006) and studies on specific domains of neology (Elsen 2004) managed to shed light on the spread of new words in more specific domains.
- The advent of **web corpora** in the last two decades has provided researchers with bigger and less formal data to study lexical innovation.
 - The sheer size of big corpora → bigger samples
 - monitoring corpora (Davies 2013): tracking dynamics of diffusion, closer to coining
 - In particular, a range of tools enabled the creation of specialized corpora for the investigation of neologisms. (Renouf, Kehoe & Banerjee 2006, Kerremans, Stegmayr & Schmid 2012, Lemnitzer 2018, Gérard 2017, Cartier 2017)
 - the nature of web corpus data is particularly suitable for investigating lexical innovations as
 - * language on the web is very creative,
 - * more informal sources, bigger spectrum of language use
 - * new words often first occur on the web
 - * and use on the web significantly influences whether these new formations catch on or not.
 - Web corpora thus promise insights into diffusion across
 - * contexts: e.g. whether new words such as *blockchain* are increasingly used in less formal
 - * cotexts: e.g. whether new words such as are increasingly used in more divers

3.2 Going beyond frequency

The conventionality of linguistic units is commonly assessed by counting how often they are found to be used in linguistic corpora, with high frequencies of occurrence seen as indicators of high levels of conventionality. Diffusion as a process that drives increasing conventionalization is thus commonly assumed to be reflected by increases in the usage frequency of linguistic innovations. Previous research on lexical innovation has been largely limited to this approach and has evaluated the spread and the overall success of new words on the basis of the number of tokens found in linguistic corpora. This paper takes usage frequency as a baseline and uses social network analysis to go beyond frequency to discover sociolinguistic dynamics of diffusion and conventionality that have eluded previous frequency-based approaches.

Frequency measures are widely used to study linguistic phenomena on all levels, from investigating phonological preferences between communities, to studying the increasing establishment of grammatical constructions like the *going to*-future over time, to assessing the degree to which words are conventional lexical units of a language. Usage frequency is thus commonly used by a diverse set of linguistic sub-disciplines. From a structural perspective, for example, co-occurrence frequencies of multi-word units such as *handsome man* are taken as an indicator for whether these are free combinations or more or less fixed collocations in a language system. Historical linguistics investigates phenomena like language change and grammaticalization, by analysing changes in usage frequency of certain constructions like the *going to*-future over time. Cognitive and psycholinguistic research commonly relies on frequency measures to approximate the degree to which speakers are familiar with words that are presented as linguistic stimuli in experiments to control for effects on experimental results.

The reliance on usage frequency as a measure for different phenomena in these diverse research contexts has faced substantial criticism. Stefanowitsch & Flach 2017 provide a good overview of the theoretical assumptions and problems that underlie frequency-based approaches in corpus linguistics.

(1) highly socially indexical and thus especially prone to be used only by certain sub-communities, (2) topical which makes freq. less reliable bc. it fails to capture ‘dormant’ passive knowledge of the words

freq. esp.
insuffi-
cient for
lex. inn.

When assessing the suitability of usage frequency as a measure for the diffusion and conventionality of neologisms a set of assumptions underlying the frequency-based approach need to be disentangled. While these theoretical and methodological considerations generally apply to all corpus-linguistic work, the focus will be on the current issue of lexical innovation.

I adopt Schmid’s EC-Model (Schmid 2020) as a framework for defining and delimitating the concepts of ‘conventionalization’ and ‘diffusion’.

[...] I define diffusion as a process that brings about a change in the number of speakers and communities who conform to a regularity of co-semiotic behaviour and a change in the types of cotexts and contexts in which they conform to it.

Aside from the sociolinguistic perspective on diffusion (‘number of speakers and communities’), Schmid conceptualizes diffusion as a multi-dimensional process that also takes into account changes on the syntagmatic (‘cotexts’) and pragmatic (‘contexts’) level. This paper will focus on the sociolinguistic dimension of diffusion and will leave an integrative approach including all three perspectives for further research.

Applying this definition to the context of lexical innovation thus implies that the successful diffusion of a new word is marked by it being known and used by an increasing ‘number of speakers and communities’.

By contrast, in a strict sense, usage frequency counts of a lexeme represent the total number of tokens produced by *all* speakers who have *contributed* to the *target text corpus*. The discrepancy between the theoretical definition of diffusion adopted here and

the exact information contained by frequency show that this operationalization relies on a number of assumptions that let it only approximate the construct to be measured.

Firstly, usage frequency does not provide direct information as to how many *individual speakers* have used a new word. Especially in the case of neology, there are certain new words that are disproportionally used and propagated by a relatively small, but more active and dedicated users of the new term. This leads to high overall frequency counts which falsely suggest that larger parts of the speech community have adopted the term.

Secondly, usage frequency only captures active uses of a term and fails to include how many speakers have been passively exposed to neologisms. In the context of entrenchment, Stefanowitsch & Flach (2017) refer to this problem as the ‘corpus-as-input’ and ‘corpus-as-output’ hypothesis. The underlying assumption is that the output of the speakers who have contributed to a corpus can serve as an approximation for the potential linguistic input of a comparable speaker group. Frequency thus reflects the ‘usage intensity’ of neologisms in the speech community which indicates the degree of entrenchment in individual speakers as well as an approximation of the conventionality of the neologisms in the speech community. In the case of lexical innovation this can be problematic as questionnaires studies on the use of neologisms (Kerremans 2015) show that many speakers report that they have come across target neologisms, but have not actively used them in discourse. Relying on frequency counts only can thus often lead to underestimating the degree of diffusion of neologisms.

Thirdly, usage frequency fails to capture where new words diffuse across ‘communities of speakers’, as suggested by Schmid’s definition. This is, of course, a consequence of the fact that frequency counts cannot provide direct information about the number of speakers involved in the diffusion of neologisms, as was pointed out in the first two points above. New words often stem and quickly spread within tight-knit communities of practice that share common attitudes or interests. Frequency measures alone cannot detect whether neologisms only show increasing usualization within these groups or whether they diffuse and become conventional in other parts of the speech community, which represents an essential feature of the sociolinguistic dimension of diffusion.

how much the words might have diffused outside the *target text corpus*
temporal dynamics (e.g. *millennium bug*)

4 Data

4.1 Neologism sample

I base my empirical study on a selection of 100 neologisms and study their use on Twitter from the early stages of their diffusion to the end of 2018. The lexemes were selected to cover a broad spectrum of lexical innovation. Previous work by Kerremans (2015: 115–147) has identified four main clusters of neologisms on the conventionalization continuum: ‘non-conventionalization’, ‘topicality or transitional conventionalization’, ‘recurrent semi-conventionalization’ and ‘advanced conventionalization’. My sample largely contains neologisms taken from the NeoCrawler, which uses dictionary-matching to retrieve a semi-automatic, bottom-up selection of recent neologisms on the web and on

Twitter (Kerremans et al. 2019). I have additionally included lexemes that were statistically identified as increasing in frequency on Twitter in recent years by Grieve, Nini & Guo (2016).

dimensions

overall degree of diffusion (synchronic): successful vs. unsuccessful: *usage frequency*, *degree centralization* no success limited advanced

temporal dynamics of diffusion (diachronic) stability: stable vs. topical *coefficient of variation* trend increasing: potential diffusion decreasing: potential centralization time window: speed and lifespan

Case studies criteria covering clusters of neologism candidates frequency counts comparable cases (Kerremans 2015) no diffusion: *microflat* limited topical: *poppygate* centralized: *alt-left* decreasing: *solopreneur* advanced diffusion: advanced: *upcycling* increasing: *hyperlocal*

4.2 Twitter data

Twitter¹ is a popular micro-blogging and social media platform that was started in 2006. one of the most popular social media platforms

representativity limited: certain user groups over-represented. good enough: generally okay (Grieve et al. 2019)

advantages of data hotbed: influential hotbed for technical, cultural and linguistic innovations longitudinal: retrospective, starting from early attestations big data social network information

scope of the present study: I only study diffusion of new words *on Twitter* considering both frequency and social network characteristics might be more robust and allow more generalization than just frequency

data retrieval I have used the tool *twint* to download all tweets containing the target word as searchable via the Advanced Search Function on Twitter's official website

data overview: table

5 Method

anatomy of a tweet

post-processing treatment of hashtags filter uses as user handles duplicate removal

social network analysis software R: *tidygraph*, *ggraph* *Gephi* constructing the networks elements: legitimate (Goel et al. 2016) users = nodes interactions (replies, mentions) = edges subsetting: time slices start of diffusion process: ≥ 3 edges 4 phases of equally sized time windows network metrics centralization network density visualizations

result: interactional networks The resulting networks are interactional rather than static. (Goel et al. 2016) This makes them more similar to communities of practice than to traditional sociolinguistic networks based on static speaker characteristics such socio-economic status. In the case of lexical innovation networks that are based on

¹<http://www.twitter.com>

Table 1: Total usage frequency in the corpus.

(a) Least frequent neologisms in the sample.		(b) Examples around the median.	
lexeme	freq	lexeme	freq
microflat	426	white fragility	26,688
dogfishing	399	monthiversary	23,607
begpacker	283	helicopter parenting	26,393
halfalogue	245	deepfake	20,101
rapugee	182	newsjacking	20,930
bediquette	164	twittosphere	20,035

(c) Most frequent examples in the sample.		(d) Case study selection.	
lexeme	freq	lexeme	freq
tweeter	7,367,174	alt-right	1012150
fleek	3,412,807	solopreneur	282026
bromance	2,662,767	hyperlocal	209937
twitterverse	1,486,873	alt-left	167124
blockchain	1,444,300	upskill	57941
smartwatch	1,106,906	poppygate	3807

whether speakers provide valuable information. In cases such as *alt-left*, for example, interactional networks show whether usage of the term remains centralized to a tight-knit community of speakers or whether it diffuses to be used by other sub-communities. Whether communities are distinct depends on whether users communicate with each other. While the reasons for these communicative affiliations remain unknown (age, gender, socio-economic status), they are certainly real in mutually engage in communicative interaction. It would be interesting to complement this information with static information, however such data are currently not obtainable.

6 Frequency

6.1 Total usage frequency

As described in Section XXX, the degree of conventionality is often approximated by a how many times a word has been used in a corpus. A common way to use this information is to rely on cumulated frequency counts which sum up the total number of uses.

The present sample of neologisms covers a broad frequency range. Table 1 presents candidates from four groups: six examples around the minimum, the median and the maximum total usage frequency as well as six words that will serve as case studies in the following sections.

In a strict sense, usage frequency only captures how many tokens of a word were

produced by all speakers who have contributed to the corpus at hand. Investigating the degree to which new words diffuse to new speakers and speaker communities on the basis of frequency counts thus depends on several inferences that are commonly accepted as sufficiently reliable.

1. Frequency counts indicate how many speakers have used the term.
2. The number of speakers who have used the term indicates how many speakers are familiar with the term, whether they have actively used it or not.
3. The number of speakers who are familiar with the term indicate how many communities of speakers are familiar with the term.

These assumptions are to a large extent plausible and have empirically been proven to be effective for investigating both degrees of entrenchment of lexemes in individual speakers as well their conventionality in the speech community.

The frequency-based division of neologisms into groups as presented in Table 1, for example, largely seems to fit common intuition. Neologisms that show very high frequency counts such as *smartwatch* have certainly been used by (1) many speakers. It is also unlikely that its roughly 1 million active uses stem exclusively from (3) one or few tight-knit communities of techno-enthusiasts and that the rest of the speech community (2) has never been exposed to the term.

However, even among the group of high-frequency lexemes in Table 2c, words that show similar total usage frequencies such as *twitterverse* and *blockchain*, for which semantic transcriptions even seem unnecessary, might indeed differ significantly regarding their conventionality in different parts of the speech community. For assessing and comparing the pathways of diffusion of less-established neologisms like *hyperlocal* and *solopreneur* total frequency counts alone provide a very limited picture.

6.2 Cumulative frequency

Cumulative frequency plots can supplement total frequency counts by additional information about the temporal dynamics of diffusion. Figure presents this information for all six cases.

Most importantly, lifespan 1 comparison: e.g. *alt-left* vs. *alt-right*
introduce cases

Potential distortions uses != users: This can distort the picture, e.g. if some speakers have a much stronger preference to use the term than the average or the amount of words contributed by by each speaker is not balanced.

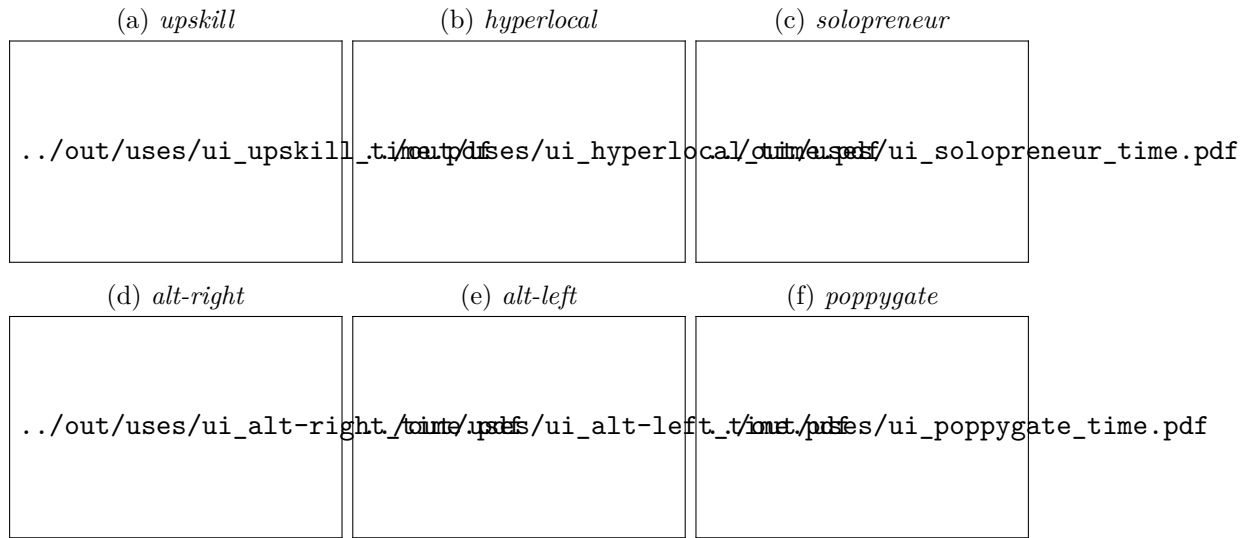
X is most frequent Y is oldest *poppygate*

²For better visibility *alt-right* was omitted from this plot because of its high usage frequency.

Figure 1: Cumulative increase in usage frequency for case studies.²

../out/freq_cum_cases.pdf

Figure 2: Temporal dynamics in usage frequency for case studies.



6.3 Temporal dynamics of usage intensity

instead of cumulative counts we now look at absolute frequency counts over time (in monthly bins)

case studies

different patterns stability trend speed of diffusion

stability: shows that freq. is problematic ‘dormant’ spikes distort representativity of frequency for degree of conventionality underestimate: *poppygate* not forgotten in troughs overestimate: cumulating hides the fact that words like *millenium* do get lost full sample

coefficient of variation most volatile least volatile

volatile patterns are the rule than the exception for *lexical innovation* due to nature of *lexical* innovation bound to cultural conceptual salience (variable ‘semantic carrying capacity’ (Nini et al. 2017)) needs to be accounted for

trend increasing: looks successful decreasing: looks unsuccessful

going beyond frequency In the following sections I will assess the value of usage frequency and compare and complement it with social network information about the diffusion of lexical innovations.

7 Social networks of diffusion

7.1 Centralization over time

going beyond frequency

def. diffusion: numbers of users communities

subsetting / time slices start of diffusion process 4 quarters

explain: degree centralization
case studies
example where freq. meets nets
example where nets add to freq.: *alt-left*

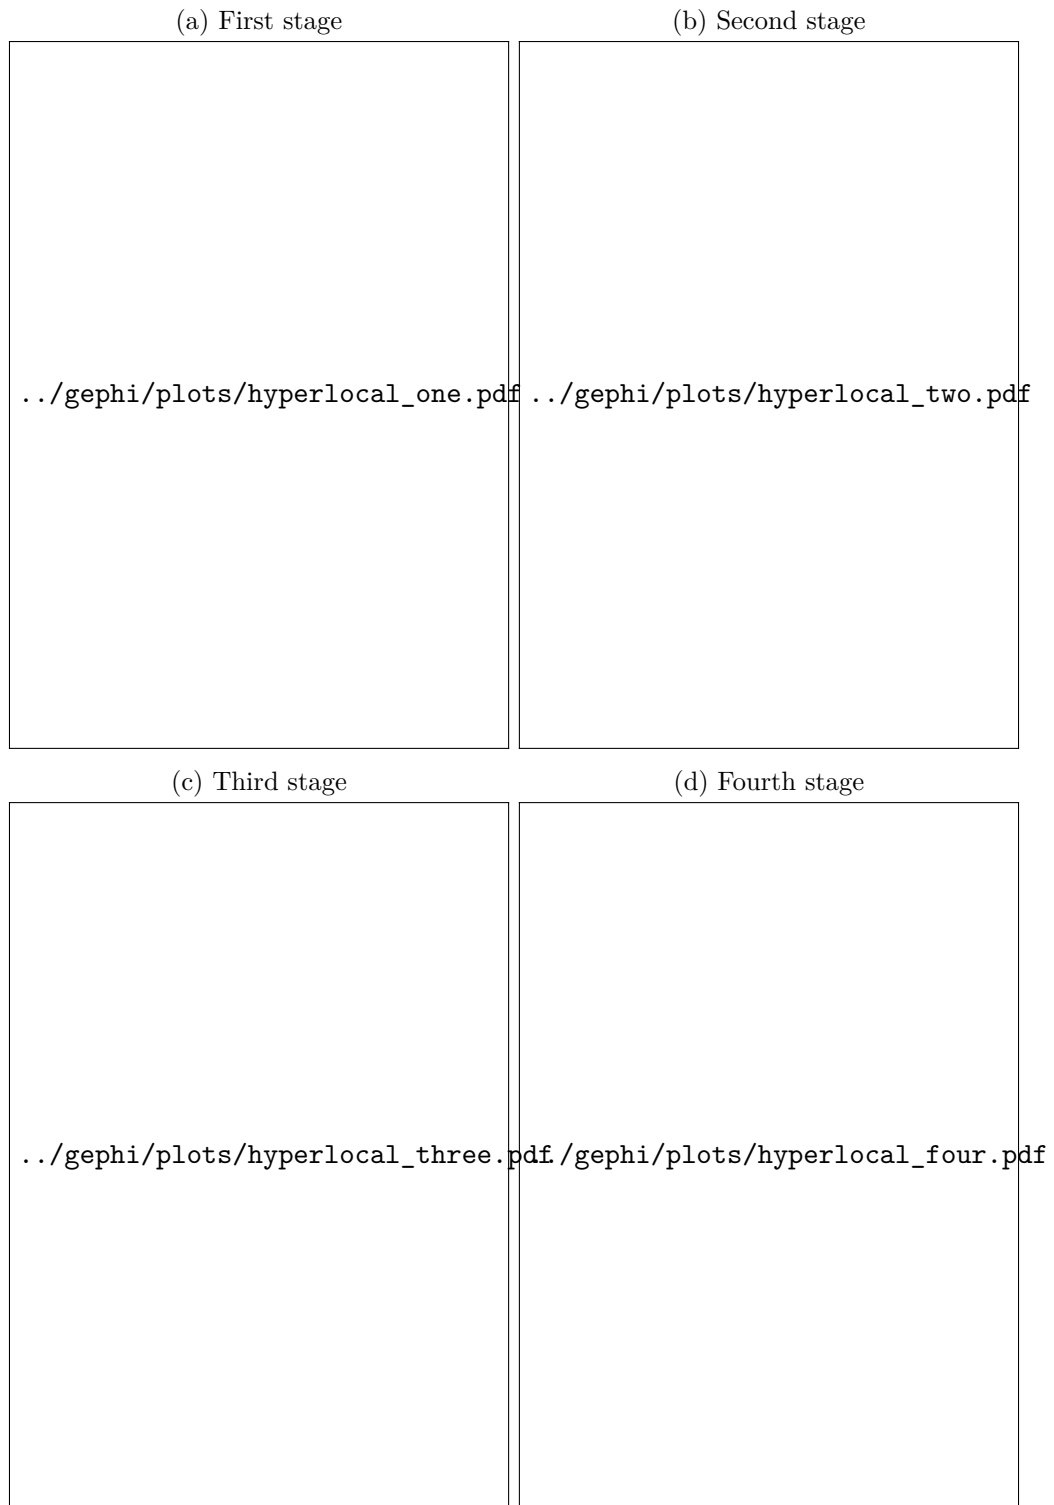
7.1.1 Overview of changes in centralization for case studies.

Figure 4: Degree centralization over time for case study words.

../out/cases_cent_diac.pdf

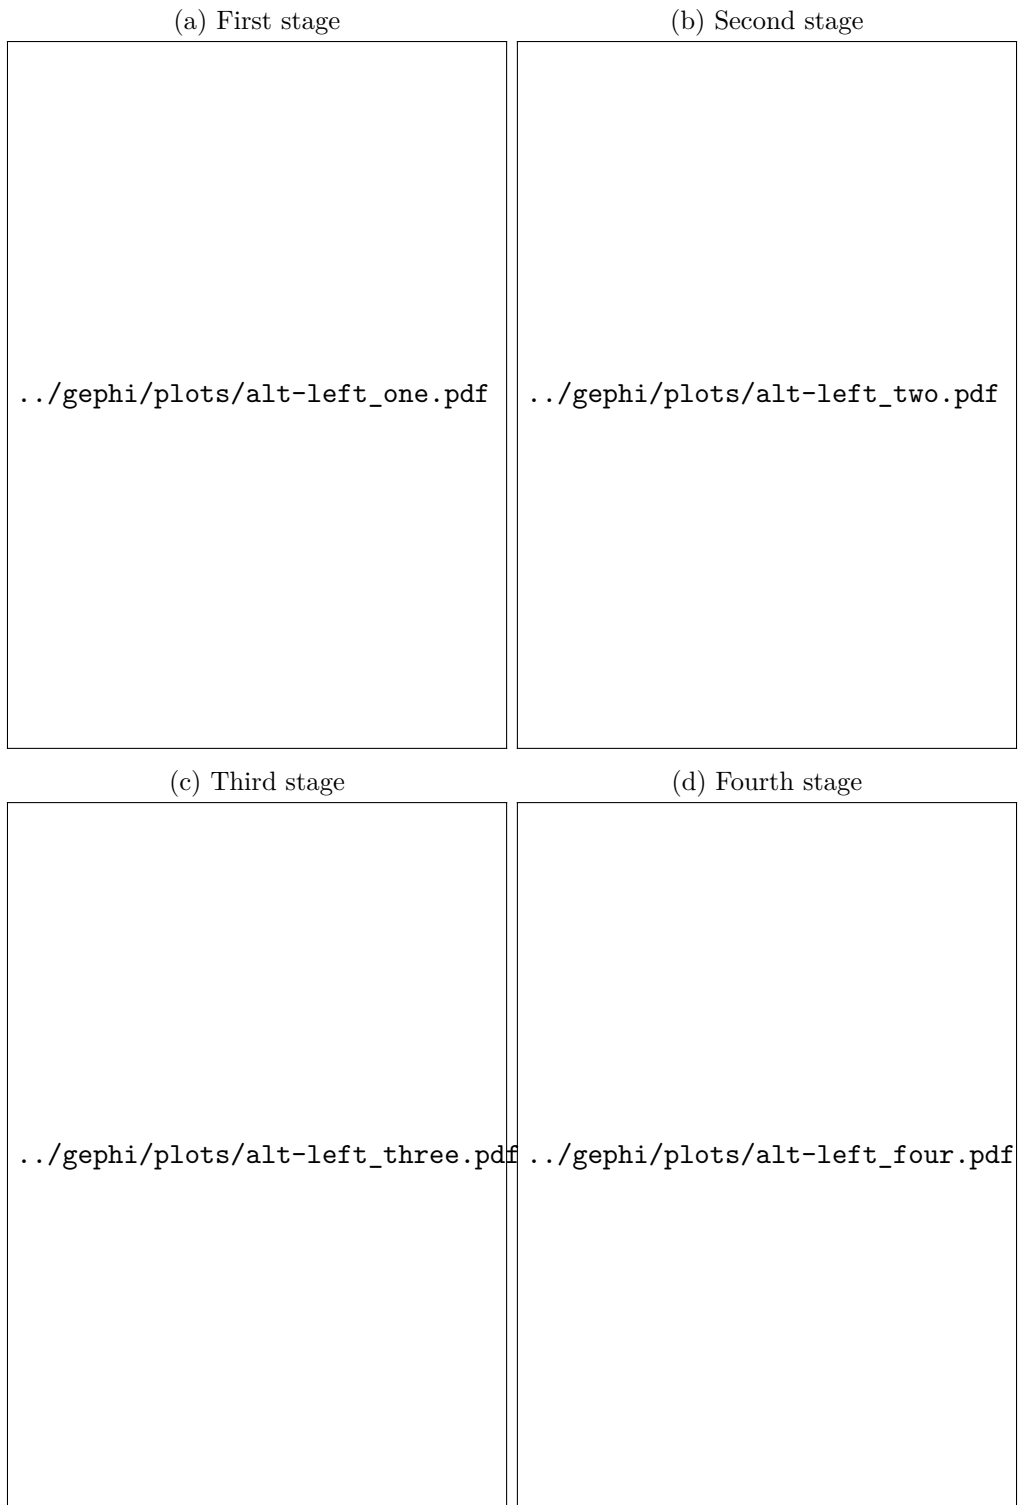
7.1.2 Advanced / increasing: *hyperlocal*

Figure 5: Social network of diffusion for *hyperlocal* over time.



7.1.3 Limited / limited: *alt-left*

Figure 7: Social network of diffusion for *alt-left* over time.



7.1.4 Full sample

density successful unsuccessful
biggest changes

7.2 Overall centralization

most diffused least diffused


8 Networks vs. frequency

8.1 Correlation

There is a significant correlation ($p = 0.015$) between frequency and centralization

8.2 Discrepancies

However, we also see discrepancies
plots



cluster analysis
freq. overestimating topical propaganda: *alt-right*, *alt-left*, *covfefe*, *birther* Brexit
terms: *Brexit*eer, *Brexit*er, *Brexit* nerds technical
freq. underestimating: XXX topical words
Social network metrics certainly add to freq.

9 Conclusion

freq. proves to be a pretty good indicators but temporal dynamics important social
network dynamics important, esp. w.r.t. new words cross-checking other data sources
(NOW corpus) shows validity social network analysis can be an important tool for soci-
olinguistics

References

- Bauer, Laurie. 1983. *English word-formation*. Cambridge university press.
- Blythe, Richard A. & William Croft. 2012. S-curves and the mechanisms of propagation in language change. *Language* 88(2). 269–304.
- Cartier, Emmanuel. 2017. Neoveille, a web platform for neologism tracking. In *Proceedings of the Software Demonstrations of the 15th Conference of the European Chapter of the Association for Computational Linguistics*, 95–98. Valencia, Spain: Association for Computational Linguistics. <http://aclweb.org/anthology/E17-3024>.
- Davies, Mark. 2013. *Corpus of News on the Web (NOW) - 3+ Billion Words from 20 Countries, Updated Every Day*. <https://www.english-corpora.org/now/>.
- Eisenstein, Jacob, Brendan O'Connor, Noah A. Smith & Eric P. Xing. 2014. Diffusion of lexical change in social media. *PLOS ONE* 9(11). 1–13. <https://doi.org/10.1371/journal.pone.0113114>.
- Elsen, Hilke. 2004. Neologismen.
- Gérard, Christophe. 2017. The logoscope: Semi-automatic tool for detecting and documenting the contexts of french new words.
- Goel, Rahul, Sandeep Soni, Naman Goyal, John Paparrizos, Hanna Wallach, Fernando Diaz & Jacob Eisenstein. 2016. The social dynamics of language change in online networks. In Emma Spiro & Yong-Yeol Ahn (eds.), *Social informatics*, 41–57. Cham: Springer International Publishing.
- Granovetter, Mark S. 1977. The strength of weak ties. In Samuel Leinhardt (ed.), *Social networks*, 347–367. Academic Press. <https://doi.org/10.1016/B978-0-12-442450-0.50025-0>.
- Grieve, Jack. 2017. Geographical patterns of lexical innovation. Workshop 'Diffusion of Lexical Innovations', LMU Munich.
- Grieve, Jack. 2018. Natural selection in the modern English lexicon. In *Proceedings of EVOLANG XII*. Torun, Poland.
- Grieve, Jack, Chris Montgomery, Andrea Nini, Akira Murakami & Diansheng Guo. 2019. Mapping lexical dialect variation in British English using Twitter. *Frontiers in Artificial Intelligence* 2. 11. <https://www.frontiersin.org/article/10.3389/frai.2019.00011>.
- Grieve, Jack, Andrea Nini & Diansheng Guo. 2016. Analyzing lexical emergence in Modern American English online. *English Language and Linguistics* (21). 99–127.
- Grieve, Jack, Andrea Nini & Diansheng Guo. 2018. Mapping lexical innovation on American social media. *Journal of English Linguistics*.
- Hébert-Dufresne, Laurent, Samuel V. Scarpino & Jean-Gabriel Young. 2020. Macroscopic patterns of interacting contagions are indistinguishable from social reinforcement. *Nature Physics*. <https://doi.org/10.1038/s41567-020-0791-2>.
- Hohenhaus, Peter. 1996. Ad-hoc-wortbildung. Terminologie, typologie und theorie kreativer wortbildung im englischen.
- Hohenhaus, Peter. 2006. Bouncebackability. A web-as-corpus-based study of a new formation, its interpretation, generalization/spread and subsequent decline. *SKASE Journal of Theoretical Linguistics* 3. 17–27.

- Kerremans, Daphné. 2015. *A Web of New Words*. Bern, Schweiz: Peter Lang. <https://doi.org/10.3726/978-3-653-04788-2>.
- Kerremans, Daphné, Jelena Prokić, Quirin Würschinger & Hans-Jörg Schmid. 2019. Using data-mining to identify and study patterns in lexical innovation on the web: The NeoCrawler. *Pragmatics and Cognition* 25(1). 174–200.
- Kerremans, Daphné, Susanne Stegmayr & Hans-Jörg Schmid. 2012. The NeoCrawler: Identifying and retrieving neologisms from the internet and monitoring ongoing change. In *Current Methods in Historical Semantics*, 59–96. Berlin: Mouton de Gruyter.
- Labov, William. 2007. Transmission and diffusion. *Language* 83(2). 344–387.
- Lemnitzer, Lothar. 2018. *Wortwarte*. <http://www.wortwarte.de/> (14 January, 2018).
- Lipka, Leonhard. 2005. Lexicalization and institutionalization: Revisited and extended. *SKASE Journal of Theoretical Linguistics* 2(2). 40–42.
- Lu, Fred Sun, Suqin Hou, Kristin Baltrusaitis, Manan Shah, Jure Leskovec, Rok Susic, Jared Hawkins, John Brownstein, Giuseppe Conidi, Julia Gunn, Josh Gray, Anna Zink & Mauricio Santillana. 2018. Accurate influenza monitoring and forecasting using novel internet data streams: A case study in the boston metropolis. *JMIR Public Health and Surveillance* 4(1). e4. <https://doi.org/10.2196/publichealth.8950>.
- Milroy, James. 1992. *Linguistic variation and change: On the historical sociolinguistics of English*. Oxford: Blackwell.
- Milroy, James & Lesley Milroy. 1985. Linguistic change, social network and speaker innovation. *Journal of Linguistics* 21(2). 339–384. <https://www.cambridge.org/core/article/linguistic-change-social-network-and-speaker-innovation1/EB30A7117CC09F6EDA5255BF9D788D5A>.
- Nevalainen, Terttu. 2015. Descriptive adequacy of the S-curve model in diachronic studies of language change. *Studies in Variation, Contacts and Change in English* 16. <http://www.helsinki.fi/varieng/series/volumes/16/nevalainen/>.
- Nini, Andrea, Carlo Corradini, Diansheng Guo & Jack Grieve. 2017. The application of growth curve modeling for the analysis of diachronic corpora. *Language Dynamics and Change* 7(1). 102–125.
- Pew Research Center. 2019. *National Politics on Twitter: Small Share of U.S. Adults Produce Majority of Tweets*. <https://www.people-press.org/2019/10/23/national-politics-on-twitter-small-share-of-u-s-adults-produce-majority-of-tweets/>.
- Renouf, Antoinette, Andrew Kehoe & Jayeeta Banerjee. 2006. WebCorp: an integrated system for web text search. *Language and Computers* 59(1). 47–67.
- Rogers, Everett M. 1962. *Diffusion of innovations*. New York: Free Press of Glencoe.
- Schmid, Hans-Jörg. 2008. New words in the mind: Concept-formation and entrenchment of neologisms. *Anglia – Zeitschrift für englische Philologie* 126. 1. <http://www.degruyter.com/view/j/angl.2008.126.issue-1/angl.2008.002/angl.2008.002.xml>.
- Schmid, Hans-Jörg. 2016. *English morphology and word-formation - An introduction*. 2nd edn. Berlin: Erich Schmidt Verlag.
- Schmid, Hans-Jörg. 2020. *The dynamics of the linguistic system. - Usage, Conventionalization, and Entrenchment*. Oxford: Oxford University Press.

- Stefanowitsch, Anatol & Susanne Flach. 2017. The corpus-based perspective on entrenchment. In Hans-Jörg Schmid (ed.), *Entrenchment and the psychology of language learning: How we reorganize and adapt linguistic knowledge*, 101–128. Boston, USA: American Psychology Association and de Gruyter Mouton.
- West, Robert, Hristo S. Paskov, Jure Leskovec & Christopher Potts. 2014. Exploiting social network structure for person-to-person sentiment analysis. *CoRR* abs/1409.2450. <http://arxiv.org/abs/1409.2450>.